

APRIL 2024



Energy-Accuracy Trade off in Hyperparameter Tuning

A study on how hyperparameter tuning affects the accuracy and energy consumption of machine learning models.

Abdulrahman Alshahrani, Abdallah Al-Sukhni

Department of Computer Science, The University of Texas at Austin

Agenda:

1. Energy consumption of Machine Learning
2. What is Hyperparameters Tuning?
3. The issue with Grid Search
4. Our Experiment
5. Our Results
6. Conclusion

Machine Learning Energy Consumption

- AI could consume **85-134** TWh by 2027 ^[1]
- Competes with **Netherlands** annual energy consumption. ^[1]
- Training GPT-3 emitted **502** metric tons of carbon, equivalent to driving **112** gasoline powered cars for a year. ^[2]

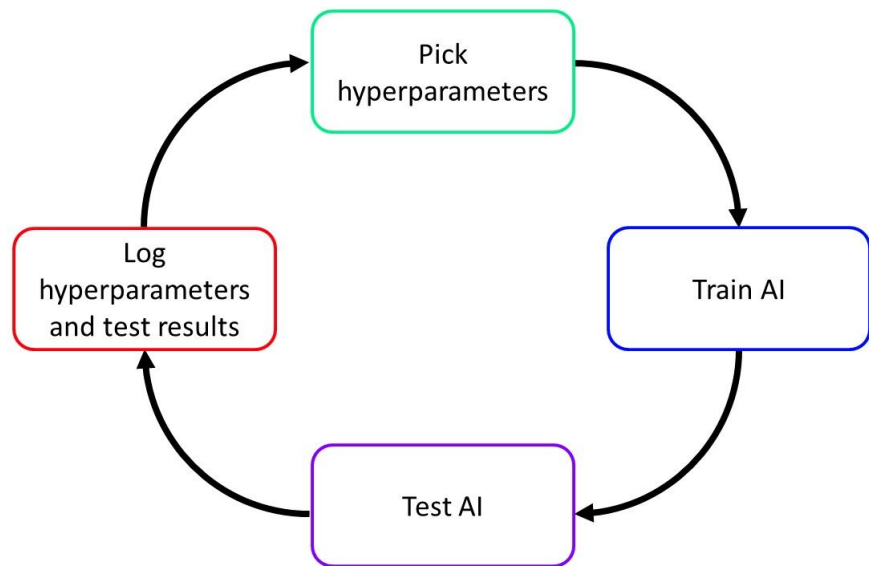
[1] How much electricity does AI consume?

Vincent

<https://www.theverge.com/24066646/ai-electricity-energy-watts-generative-consumption>

[2] <https://news.climate.columbia.edu/2023/06/09/ais-growing-carbon-footprint/>

How does an ML Model Works







But What's a Hyperparameter?

- It specifies the **structure** of your ML model.
- You set it **before** training your model.


But What's a Hyperparameter?

E.g.

- K-nearest neighbors

 ==  or  ==  ?


What K value to use?



But What's a Hyperparameter?

E.g.

- Neural Networks



How many nodes in each layer?

how many hidden layers?

How to choose the **optimal** Hyperparameters?

(aka Hyperparameter Tuning)

- First... What is “*optimal*”?
- Optimal = the highest accuracy for the model



How to choose the **optimal** Hyperparameters?

(aka Hyperparameter Tuning)



- Background experience (too hard, not reliable)

Or...

- Let the machine try many values


What values to try?

- **Testing all values**
 - **aka Grid Search**
- **Test Random values**
 - **aka Randomized Search**



The issue with Grid Search

- Exponential



- Maybe inefficient?

The issue with Grid Search

- **~2000x** increase in CO₂ [1]

Training one model (GPU)	CO ₂ e (lbs)
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468

- **77%** decrease in energy consumption vs. **1%** decrease in accuracy [2]

[1] E. Strubell, A. Ganesh, and A. McCallum, "Energy and Policy Considerations for Deep Learning in NLP," 2019.

[2] A. Brownlee, J. Adair, S. Haraldsson, and J. Jabbo, "Exploring the Accuracy -Energy Trade-off in Machine Learning."


So Who Cares?






What Others Did

- Brownlee, Adair, Haraldsson, Jabbo
- They investigated the energy and accuracy trade-offs of using grid search on multiple hyperparameter options across several datasets
- Findings revealed a clear opportunity to trade off small accuracy in exchange for a much higher reduction in energy consumption



What Others Did

- Energy savings of **~40%** with minimal reductions in classification accuracy.
- They only focused on MultiLayer Perceptrons (MLP)



What We're Doing

- Can we get similar savings with other algorithms?
 - Multilayer Perceptrons (MLPs)
 - K Nearest Neighbors
 - Decision Trees
 - Random Forest

Our Setup

- The Algorithms
- The Datasets
- Hyperparameter Grids
- The Machines
- Measuring Energy Consumption

The Algorithms

- Grid Search
 - GridSearchCV()
- Classifiers
 - DecisionTreeClassifier()
 - KNeighborsClassifier()
 - MLPClassifier()
 - RandomForestClassifier()



The Datasets

- Olivetti Faces (from scikit-learn)
 - ~400 instances
- Census Income
 - ~50k instances
- CDC Diabetes Health Indicators
 - ~ 250k instances
- RT-IoT2022
 - ~120k instances



Hyperparameter Grids

For example, MLP's:

```
param_grids = [  
    {'hidden_layer_sizes': [(100,), (50, 50)]},  
    {'hidden_layer_sizes': [(100,), (50, 50), (100, 100), (200, 200)]},  
    {'hidden_layer_sizes': [(100,), (50, 50), (100, 100), (200, 200), (300, 300)]},  
    {'hidden_layer_sizes': [(100,), (50, 50), (100, 100), (200, 200), (300, 300)],  
    'alpha': [0.0001, 0.001],  
    'activation': ['relu', 'tanh', 'logistic']}  
]
```

The Machines

- Abdulrahman's Machine:

- **CPU:** Intel Core i7-9750H, 6 cores, (Base: 2.60GHz, Max: 4.50 GHz)
- **GPU:** NVIDIA GeForce GTX 1660 Ti
- **Memory:** 32 GB

- Abdallah's Machine:

- **CPU:** Intel Core i5-9400, 6 cores, (Base: 2.90GHz, Max: 4.10 GHz)
- **GPU:** NVIDIA GeForce GTX 1660 Ti
- **Memory:** 16 GB



Measuring Energy Consumption

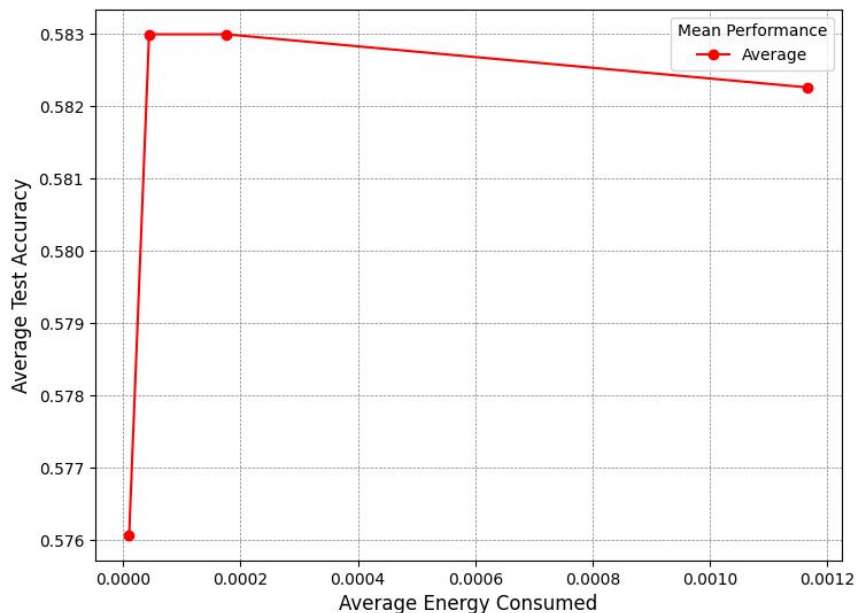
- **CodeCarbon** Python Library
 - Open Source
 - Lightweight
- Polling components for power consumption every some interval
 - We used every 5 seconds

Measuring Energy Consumption


- GPU Power
 - Using **NVIDIA Management Library** on pyNVML
- CPU Power
 - Using **Intel Power Gadget**

The Results - Decision Trees

DecisionTree - census

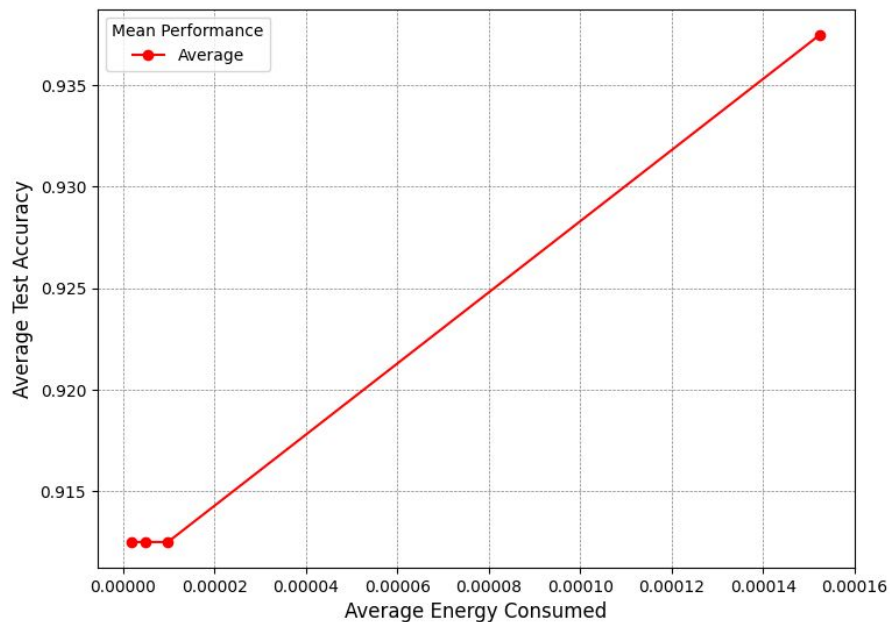


DecisionTree - CDCHealthIndicator

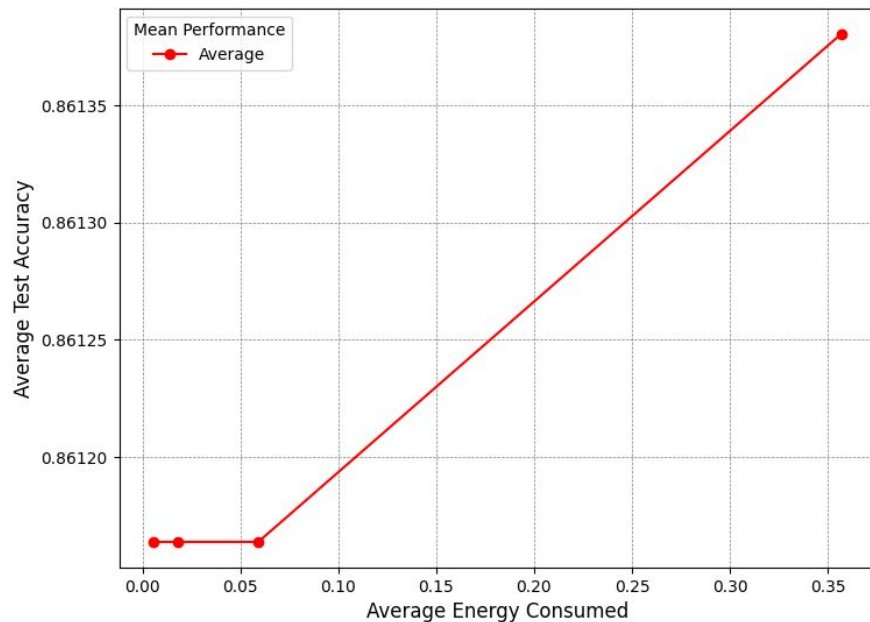


The Results - KNN

KNeighbors - olivetti

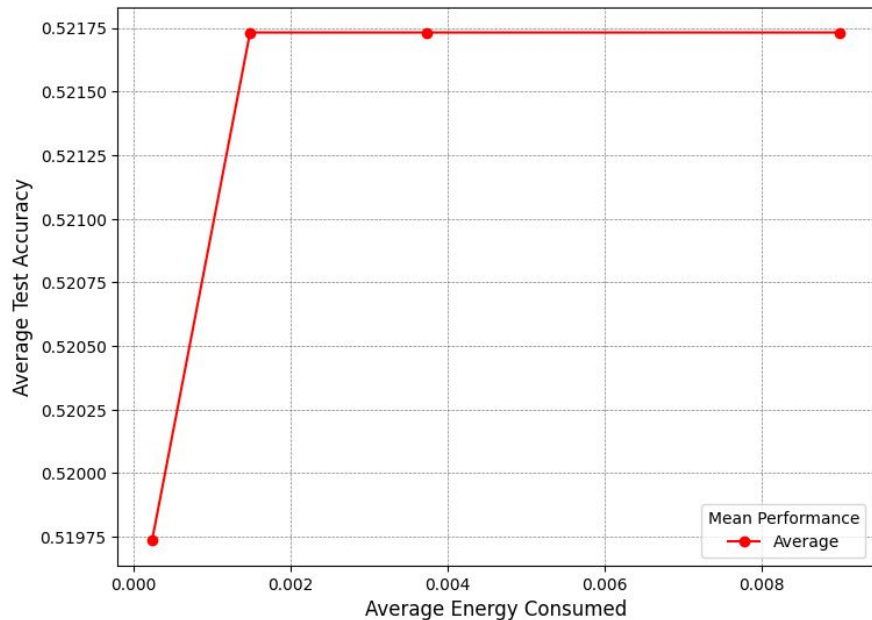


KNeighbors - CDCHealthIndicator




The Results - MLPs

MLP - census




MLP - RT-IoT2022

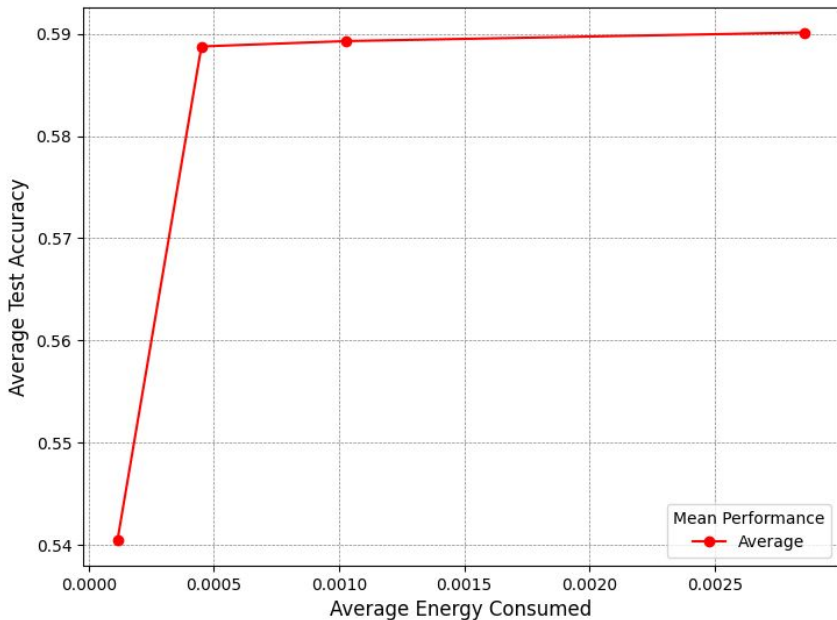


The Results - Random Forest

RandomForest - RT-IoT2022



RandomForest - census



Findings

- Most algorithms showed **minimal increase** in accuracy
- Testing additional hyperparameters increased energy consumption dramatically.
- Average increase in energy per additional grid search is **309.64%**
- Average increase in accuracy per additional grid search is **0.01%**

Conclusion

- Can we say Grid Search is inefficient?
 - Mostly yes!
- There are clear points of diminishing/no returns
- Unless
 - any increase in accuracy vital to you, Or
 - you maybe don't care about energy

Experiment Code

- Code is available at our repository:
github.com/abdomash/ml-energy-consumption
- You can run the experiment on your device
- Setup guide in the README file

